


# Web Server and Application Server

## 1) Install nginx and run it on port number 81.

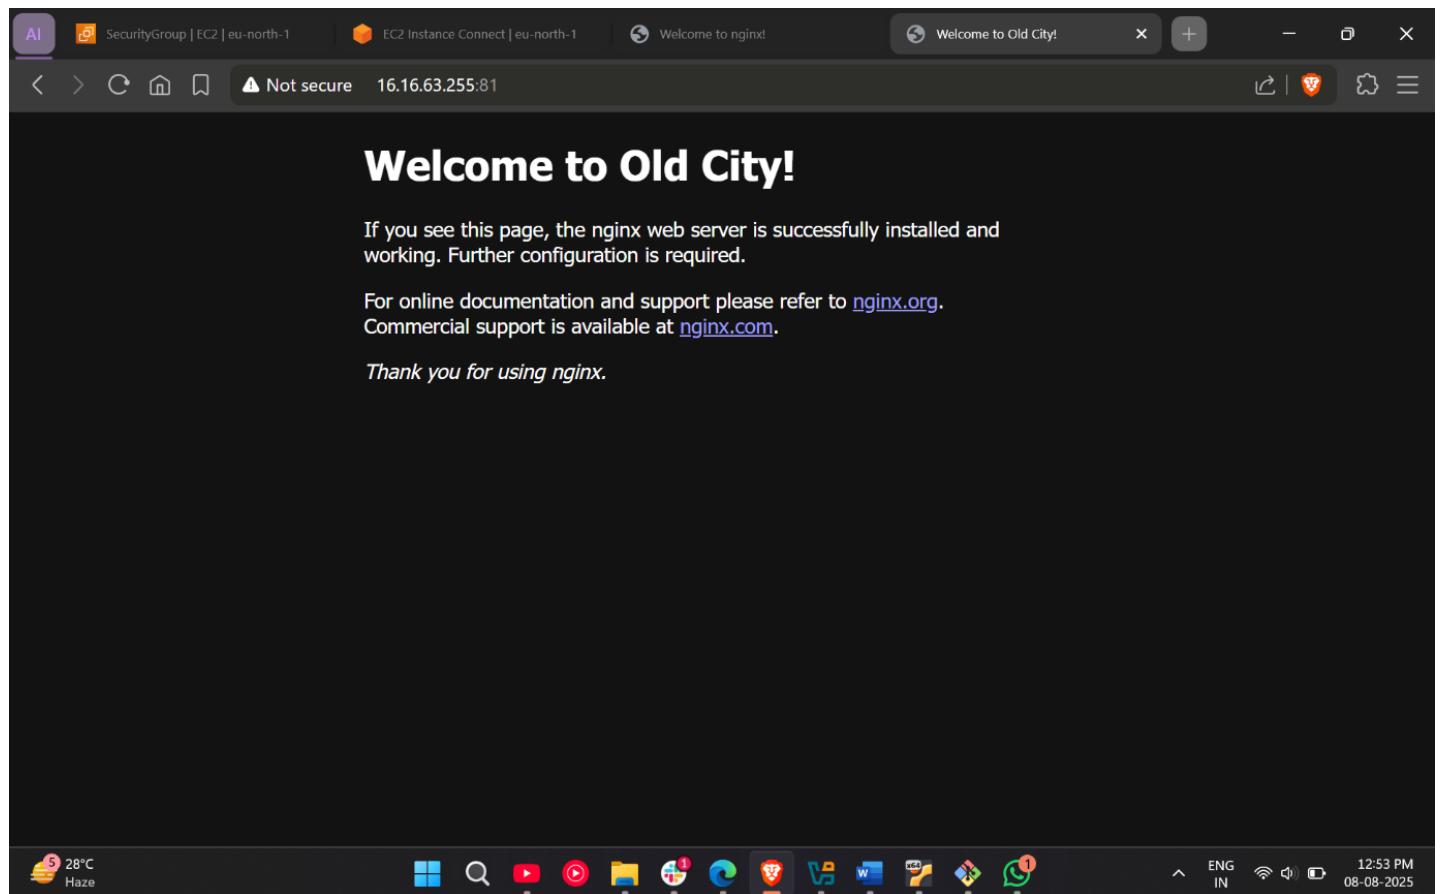


Steps to install the nginx web server:

- 1) Launch an instance on AWS and select any Amazon Linux OS
- 2) Update and upgrade the machine before installing nginx. To do that, use:  
`# yum update & # yum upgrade`
- 3) Install nginx using this command:  
`# yum install nginx -y`
- 4) Next step is to start nginx and enable it. Do that using:  
`# systemctl start nginx`  
`# systemctl enable nginx`
- 5) The next step is to make a copy of the conf file. The conf file is located at `/etc/nginx/nginx.conf`. It is always best practice to make a backup of the file in case the original file gets corrupted.  
`# cp /etc/nginx/nginx.conf /etc/nginx/nginx.conf.backup`
- 6) Edit the “nginx.conf” file and change the port number to 81 from 80.  
`# vi nginx.conf` or `vim nginx.conf` or `nano nginx.conf`  
While editing, we look for the port we need to change. It will be in the line:  

```
server {  
    listen    80;  
    listen    [::]:80;  
    server_name _;
```
- 7) Change the port to 81. Then restart nginx with:  
`# systemctl stop nginx` or `systemctl restart nginx`
- 8) Check if the configuration is done right with:  
`# nginx -t`
- 9) Copy the public IP from the machine and open a new tab in the browser. Paste the IP with the port number and press Enter. The nginx server will successfully run on port 81. Eg: 16.16.63.255:81


## 2) Deploy a sample index.html file on nginx.



Steps to deploy a sample HTML file on nginx:

- 1) Navigate to “/usr/share/nginx/html” directory and look for an index.html file.
- 2) Create a backup of the index file for safety  
`# cp index.html index.html.backup`
- 3) Open the index.html file and make the desired changes.  
`# vi index.html`
- 4) Look for the head and the body where you need to make changes, which appear in the output. Changes made in the head can be viewed on the browser’s tab title, and changes made in the body can be noticed on the main heading of the page.
- 5) Once changes are made, save and exit the HTML file.
- 6) Check the changes by opening a new tab and pasting the public IP on port 81.  
Eg: 51.21.243.49:81

### 3) Install Apache HTTPD and run it on port number 82




**It works!**



Steps to install Apache HTTPD and run it:

- 1) Update and upgrade the machine.  
# yum update  
# yum upgrade
- 2) Install HTTPD.  
# yum install httpd -y  
(Use -y to say yes before installation, so you don't have to do it later)
- 3) Start and enable HTTPD  
# systemctl start httpd  
# systemctl enable httpd
- 4) Navigate to - /etc/httpd/conf/ and look for httpd.conf file. Make a copy for backup in case the original file gets corrupted.  
# cd /etc/httpd/conf  
# cp httpd.conf httpd.conf.backup
- 5) Open the file and change the port from 80 to 82. Save the file and provide the ownership to the httpd user.  
# vi httpd.conf – Look for port 80 and change it to 82  
(To easily find the keyword, type: esc -> :/+keyword to search. Eg: esc -> :/80)  
# chown -R httpd:httpd httpd.conf
- 6) Restart HTTPD and start again. Copy the public IP and open it on port 82.  
# systemctl restart httpd (or) #systemctl stop httpd  
# systemctl start httpd
- 7) Check for the results.

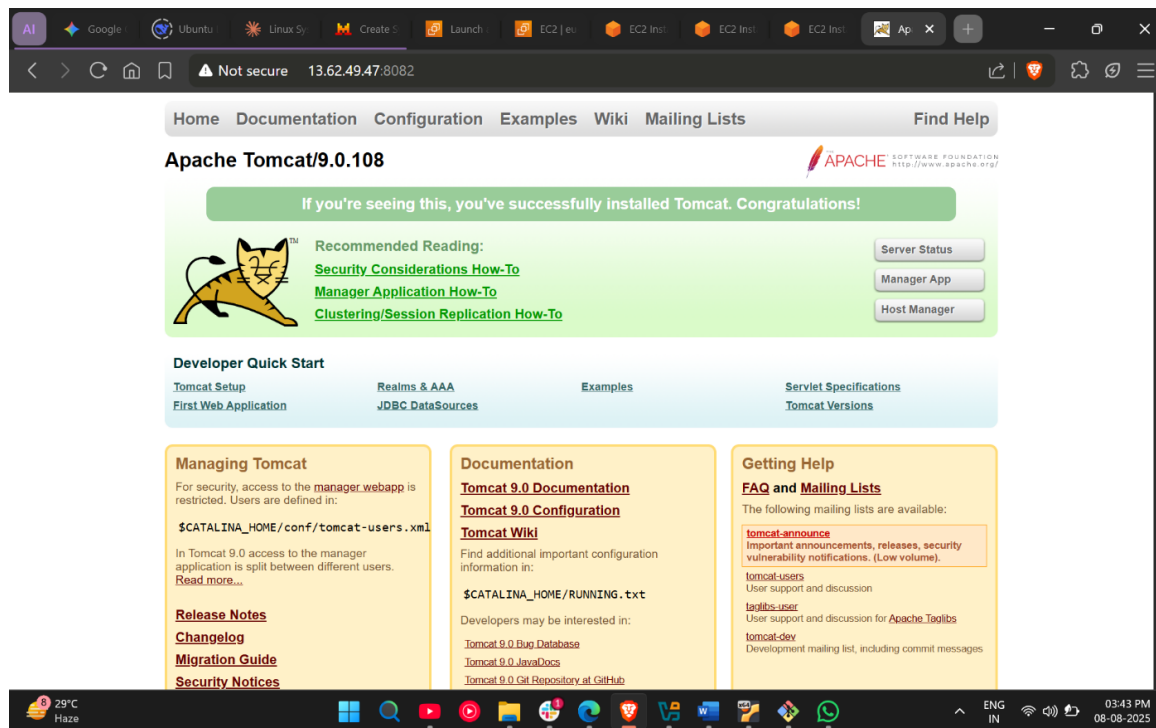
#### 4) Deploy a sample index.html file on Apache HTTPD.



Steps to edit index.html file:

- 1) Find the index.html file using either the find command or navigate to /usr/share/httpd/noindex directory.
- 2) Give appropriate permissions and ownership to the /usr/share/httpd directory to the apache user.  
# chown -R apache: /usr/share/httpd/noindex or # chown -R apache: /var/www/html/index.html
- 3) In case of /var/www/html, create a file with the name index.html and add any text to check the output.
- 4) Open the index.html file with any editor.  
# vi index.html
- 5) Make the desired changes to the file and save it.
- 6) Restart the https service. If stop command is use, make sure to start the httpd service again.  
# systemctl restart httpd or # systemctl stop httpd  
# systemctl start httpd
- 7) Open a browser and enter the public IP with the port number 82.
- 8) Check the deployed index.html file.

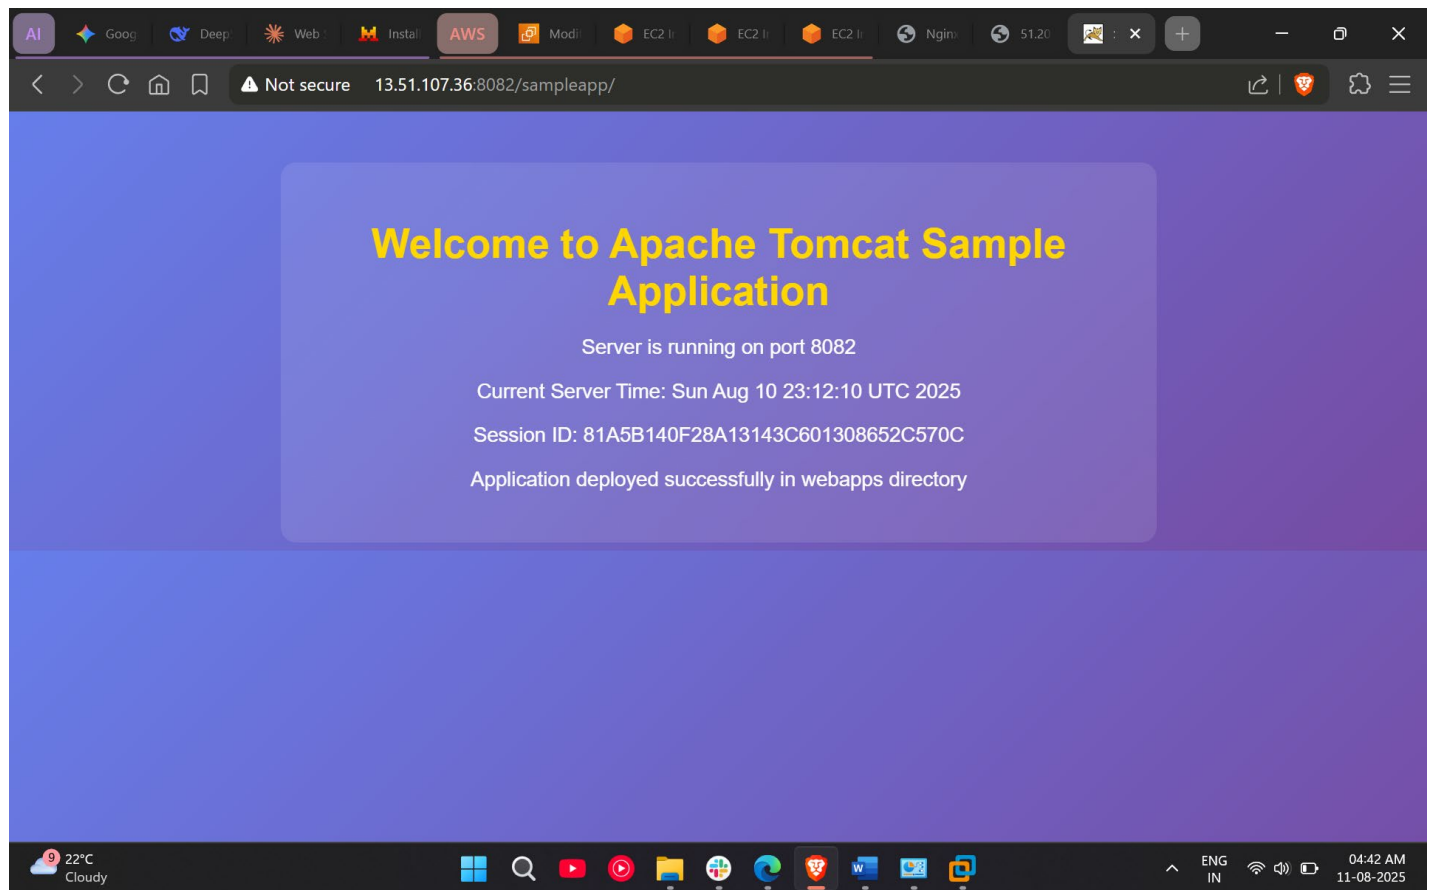
## 5) Install Apache Tomcat on port number 8082



Steps to install and run Apache Tomcat:

- 1) Create an instance and name it tomcat. Select All traffic and use source as: 0.0.0.0/0.
- 2) Update and upgrade the system.
- 3) Install the appropriate version of Java into the system.  
`# dnf install openjdk-21-amazon-corretto-devel -y`
- 4) Add a user with the name tomcat and make it a system user with a home directory in /opt/tomcat and give it a /bin/false shell. This is a good practice that helps protect the server even after a cyberattack.  
`# useradd -r -m -U -d /opt/tomcat -s /bin/false tomcat`  
r – system user  
m – to create a home directory  
U – to create a group with the same username  
d – path/to/directory  
s – providing shell, which in this case is /bin/false
- 5) Download Tomcat from the internet into the “/tmp” directory. Click on the tar file and copy the link. To download, use:  
`# wget <link_address>`
- 6) Extract the file in the “/opt/tomcat” directory.  
`# tar xvf <file_name> -C /opt/tomcat --strip-components=1`. To complete the name of the file or directory, press the **Tab** key.
- 7) Change the ownership of the directory “/opt/tomcat”, and edit the “server.xml” file. Change the port number from the default 8080 to 8082. Make a backup before editing.  
`# chown -R tomcat: /opt/tomcat`  
`# cp /opt/tomcat/conf/server.xml /opt/tomcat/conf/server.xml.backup`  
`# vi server.xml` -> Change port number from 8080 to 8082 wherever needed.
- 8) Restart Tomcat and start again from the bin directory.  
`# ./shutdown.sh && # ./startup.sh`
- 9) Open a new tab in the browser and check the output with the public IP on port 8082

## 6) Deploy a sample app on webapps



### Steps:

- 1) Create a new directory with the name "sampleapp" in the "/opt/tomcat/webapps" directory and add two files with the names "web.xml" where we will define how the web app will look like, and another file "index.jsp" where we add the content which appears of the web app.  
# mkdir /opt/tomcat/webapps/sampleapp  
# vi web.xml -> Define the web app  
# vi index.jsp -> Content of the web app
- 2) Give the ownership of the directory to tomcat  
# chown -R tomcat: /opt/tomcat/webapps/sampleapp
- 3) Restart Tomcat to deploy the application. Go to the "/opt/tomcat/conf" directory and use the commands:  
# ./shutdown.sh (or) # sh shutdown.sh  
# ./startup.sh (or) sh startup.sh
- 4) Check if the deployment is working with curl command  
# curl <http://localhost:8082/sampleapp/>
- 5) Open a new tab and copy the public IP from the instance and open it on port "8082/sampleapp/" to check the deployment.  
[http:// 13.51.107.36:8082/sampleapp](http://13.51.107.36:8082/sampleapp)

## 7) Create a “tomcat.service” file for Tomcat.

```
[root@ip-172-31-31-124 ~]# systemctl status tomcat
● tomcat.service - Apache Tomcat Web Application Container
   Loaded: loaded (/etc/systemd/system/tomcat.service; enabled; preset: disabled)
   Active: active (running) since Mon 2025-08-11 19:35:28 UTC; 58s ago
     Process: 1427 ExecStart=/opt/tomcat/bin/startup.sh (code=exited, status=0/SUCCESS)
    Main PID: 1436 (java)
      Tasks: 30 (limit: 1057)
     Memory: 180.7M
        CPU: 4.627s
    CGroup: /system.slice/tomcat.service
            └─1436 /usr/lib/jvm/java-21-amazon-corretto.x86_64/bin/java -Djava.util.logging.config.file=/opt/tomcat/conf/logging.properties -Djava.u

Aug 11 19:35:28 ip-172-31-31-124.eu-north-1.compute.internal systemd[1]: Starting tomcat.service - Apache Tomcat Web Application Container...
Aug 11 19:35:28 ip-172-31-31-124.eu-north-1.compute.internal startup.sh[1427]: Tomcat started.
Aug 11 19:35:28 ip-172-31-31-124.eu-north-1.compute.internal systemd[1]: Started tomcat.service - Apache Tomcat Web Application Container.

[root@ip-172-31-31-124 ~]# curl http://localhost:8082

<!DOCTYPE html>
<html lang="en">
  <head>
    <meta charset="UTF-8" />
    <title>Apache Tomcat/11.0.10</title>
    <link href="favicon.ico" rel="icon" type="image/x-icon" />
    <link href="tomcat.css" rel="stylesheet" type="text/css" />
  </head>
```

Steps:

- 1) Create a “tomcat.service” file in the directory - /etc/systemd/system/tomcat.service.
- 2) Edit the file using vi command and enter the following content:

[Unit]

Description=Apache Tomcat Web Application Container

After=network.target

[Service]

Type=forking

Environment=JAVA\_HOME=/usr/lib/jvm/"Check and enter the correct file/java version"

Environment=CATALINA\_PID=/opt/tomcat/temp/tomcat.pid

Environment=CATALINA\_HOME=/opt/tomcat #make sure to check the directory name and enter as it is

Environment=CATALINA\_BASE=/opt/tomcat #same as above

Environment='CATALINA\_OPTS=-Xms256M -Xmx512M -server -XX:+UseParallelGC'

Environment='JAVA\_OPTS=-Djava.awt.headless=true -Djava.security.egd=file:/dev/./urandom'

ExecStart=/opt/tomcat/bin/startup.sh

ExecStop=/opt/tomcat/bin/shutdown.sh

User=tomcat

Group=tomcat

UMask=0007

RestartSec=10

Restart=always

[Install]

WantedBy=multi-user.target

- 3) Save the file and go back to the directory “/opt/tomcat/bin” and manually shut down Tomcat.
- 4) Reload systemd and start the service. Use the following commands:
  - # systemctl daemon-reload – Reload systemd configuration files
  - # systemctl start tomcat – Start the Apache Tomcat service
  - # systemctl enable tomcat – Enable automatic startup on boot
- 5) Verify if the service is running properly.
  - # sudo systemctl status tomcat – Shows the status of the service
  - # curl <http://localhost:8082> – Shows whether the service is opening with the IP and Port number.

## 9) Configure HA Proxy server

```
[root@ip-172-31-35-136 haproxy]# ls
conf.d  haproxy.cfg  haproxy.cfg.backup
[root@ip-172-31-35-136 haproxy]# haproxy -f /etc/haproxy/haproxy.cfg -c
Configuration file is valid
[root@ip-172-31-35-136 haproxy]# systemctl start haproxy
[root@ip-172-31-35-136 haproxy]# systemctl enable haproxy
Created symlink /etc/systemd/system/multi-user.target.wants/haproxy.service → /usr/lib/systemd/system/haproxy.service.
[root@ip-172-31-35-136 haproxy]# systemctl status ha
halt.target
haproxy.service
[root@ip-172-31-35-136 haproxy]# systemctl status haproxy
● haproxy.service - HAProxy Load Balancer
   Loaded: loaded (/usr/lib/systemd/system/haproxy.service; enabled; preset: disabled)
   Active: active (running) since Mon 2025-08-11 20:27:31 UTC; 1min 19s ago
     Main PID: 26642 (haproxy)
       Status: "Ready."
        Tasks: 3 (limit: 1057)
       Memory: 6.4M
          CPU: 60ms
      CGroup: /system.slice/haproxy.service
              └─26642 /usr/sbin/haproxy -Ws -f /etc/haproxy/haproxy.cfg -f /etc/haproxy/conf.d -p /run/haproxy.pid
                └─26648 /usr/sbin/haproxy -Ws -f /etc/haproxy/haproxy.cfg -f /etc/haproxy/conf.d -p /run/haproxy.pid

Aug 11 20:27:31 ip-172-31-35-136.eu-north-1.compute.internal systemd[1]: Starting haproxy.service - HAProxy Load Balancer...
Aug 11 20:27:31 ip-172-31-35-136.eu-north-1.compute.internal haproxy[26642]: [NOTICE] (26642) : haproxy version is 2.8.3-86e043a
Aug 11 20:27:31 ip-172-31-35-136.eu-north-1.compute.internal haproxy[26642]: [NOTICE] (26642) : path to executable is /usr/sbin/haproxy
Aug 11 20:27:31 ip-172-31-35-136.eu-north-1.compute.internal haproxy[26642]: [ALERT] (26642) : config : parsing [/etc/haproxy/haproxy.cfg:29] : 'p
Aug 11 20:27:31 ip-172-31-35-136.eu-north-1.compute.internal haproxy[26642]: [NOTICE] (26642) : New worker (26648) forked
Aug 11 20:27:31 ip-172-31-35-136.eu-north-1.compute.internal systemd[1]: Started haproxy.service - HAProxy Load Balancer.
```

Steps to install and run HA-Proxy:

- 1) Create a new instance and install HA-Proxy
  - # yum install haproxy -y
- 2) Make a copy of the configuration file and then edit it. The configuration file is located in the directory “/etc/haproxy/haproxy.cfg”
- 3) Add nginx, httpd, and tomcat IPs with port numbers
  - # vi /etc/haproxy/haproxy.cfg – #Add the below text
  - global
  - log 127.0.0.1:514 local0
  - chroot /var/lib/haproxy
  - stats socket /run/haproxy/admin.sock mode 660 level admin
  - stats timeout 30s



```
user    haproxy
group   haproxy
daemon
```

#### defaults

```
mode            http
log             global
option          httplog
option          dontlognull
option          http-server-close
option          forwardfor    except 127.0.0.0/8
option          redispatch
retries         3
timeout http-request 10s
timeout queue   1m
timeout connect 10s
timeout client  1m
timeout server  1m
timeout http-keep-alive 10s
timeout check   10s
maxconn        3000
```

#### # HAProxy Statistics

##### frontend stats

```
bind *:8404
stats enable
stats uri /stats
stats refresh 30s
stats admin if TRUE
```

#### # Frontend for web servers (Nginx and Apache)

##### frontend web\_frontend

```
bind *:80
default_backend web_servers
```

#### # Backend for web servers

##### backend web\_servers

```
balance roundrobin
option httpchk GET /
server nginx_server 127.0.0.1:81 check
server apache_server 127.0.0.1:82 check
```

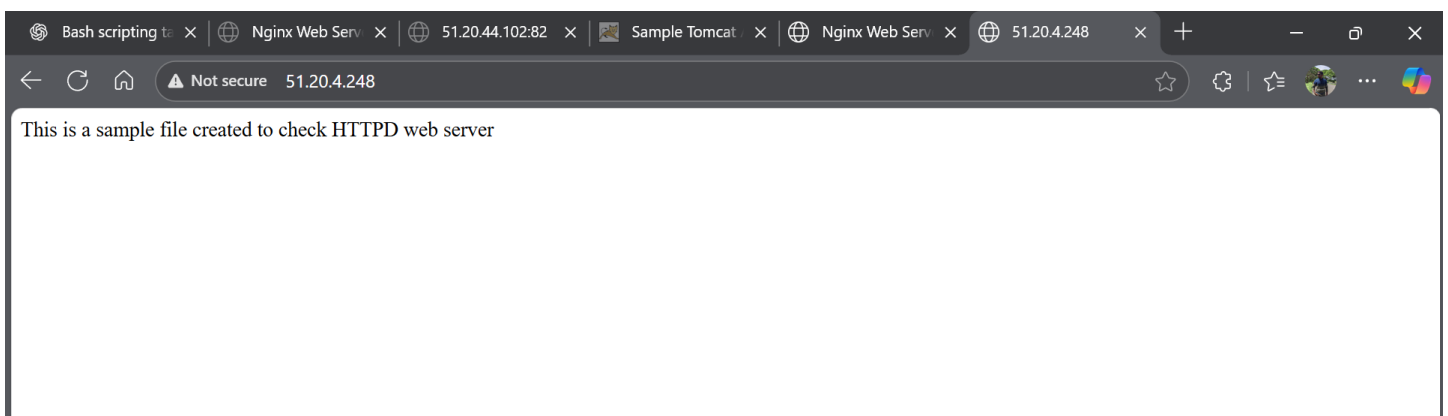
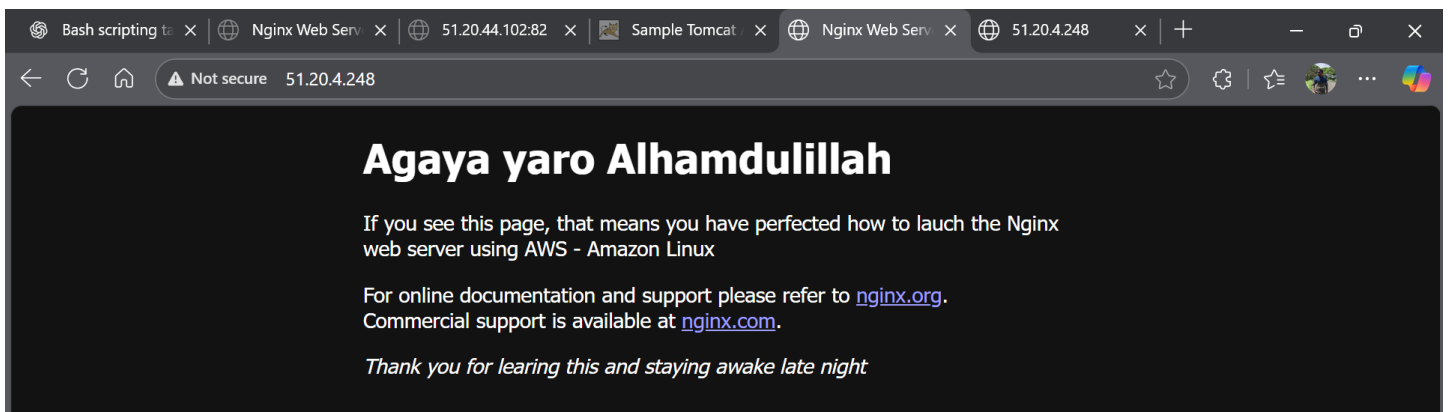
#### # Frontend for Tomcat applications

##### frontend tomcat\_frontend

```
bind *:8080
default_backend tomcat_servers
```

```
# Backend for Tomcat servers
backend tomcat_servers
    balance roundrobin
    option httpchk GET /
    server tomcat_server 127.0.0.1:8082 check
```

- 4) Test HA-Proxy configuration using:  
# haproxy -f /etc/haproxy/haproxy.cfg -c
- 5) Start and Enable HA-Proxy  
# sudo systemctl start haproxy  
# sudo systemctl enable haproxy
- 6) Verify the status of the service using:  
sudo systemctl status haproxy
- 7) On the “haproxy” server, open the file “hosts” which is located in the “/etc” directory and add the public IPs of “nginx”, “httpd”, and “tomcat”.  
# vi /etc/hosts – Add the following IPs below the already existing data  
13.60.224.33 nginx  
51.20.105.32 httpd  
51.20.83.104 tomcat
- 8) On the “nginx” server and “httpd” server, open the “hosts” file and add the “haproxy” server’s public IP, as in the above step.
- 9) Open a two new tabs and copy paste the public IP of “haproxy” to check the result of load balancer.



*Shaik Omer Farooq*  
*Batch 14*